

AURA

Agentic Understanding of Room Attention

A multimodal companion that watches what a room does — gaze, speech, typing and ink — and tells the team what matters.

USE CASE

Co-located collaborative sessions

INPUTS

Camera · Mic · Laptops · Tablet

OUTPUTS

Presenter HUD + enriched deck
(PPT-B)

TEAM

Team 2488

Feedback arrives weeks too late

No tool today lets a room full of people work on one document while a companion watches every input — attention, speech, typing and ink — and tells the team what is important.

Presenters fly blind

They learn they lost the room weeks later, through surveys.

Inputs are perishable


Typed questions and annotations vanish when the meeting ends.

Convergence needs talk

And most people can't talk and listen at the same time.

One deck, four modalities, four agents

No pixels leave the device.
Perception is local CPU.



• MODEL INSIGHTS

One mid-size LLM serves the whole room

2.6 s

• dip detection

0.4 s

• question surfacing

<4 s

end-to-end target

0

• false positives

≤700

tokens / room state

MODELS

Qwen2.5-7B-Instruct— agent reasoning · vLLM on ROCm

faster-whisper (int8, CPU)— ASR, keeps the GPU free

MediaPipe FaceMesh + PnP— attention, CPU, no CUDA

RIGHT-SIZING

~16 GB weights + 4–8 GB KV → fits a 24 GB-class GPU

Perception runs on CPU. The MI300X's 192 GB is roadmap headroom — Qwen2.5-VL screen understanding on the same server — not a requirement.

Scenario A ≈ 2.3K + 0.3K tok (2 calls) · Scenario B ≈ 10–15K tok (9-min session) · slide match 5/5 @ 0.9904 / 05
• measured today

A companion, not a dashboard

WHAT CHANGES

Feedback latency: weeks → seconds

Zero lost inputs. Every question is surfaced live or logged as unresolved in PPT-B — institutional memory by default.

Per-slide analytics. Engagement data so content improves between sessions.

Privacy by architecture. Anonymous tracklets, no pixels to the LLM — deployable where cameras are sensitive.

WATCH FOR • THE DEMO

- 01 Engagement collapses → the Analyst diagnoses the cause → the Coach nudges → the presenter recovers on camera.
- 02 Typed questions sit pending → surfaced exactly at the presenter's pause → answered from the HUD.
- 03 A calm slide where nothing fires — correct silence is the harder half of detection.
- 04 Debrief → PPT-B names the weakest slide and flags one question unresolved.